

Name: 황지영

Student ID: 2022211222

1. Implementation Results

1.1 Performance Results for Different Image Sizes


The table below shows the speedup between the baseline implementation and the optimized implementation across different image sizes. The baseline implementation time is assumed to be 100 ms. The optimized implementation time is calculated using the speedup ratio as follows:

Optimized Implementation Time (ms) = Baseline Implementation Time (ms) / Speedup

Image Size	Baseline Implementation Time (ms)	Optimized Implementation Time (ms)	Speedup
1024x1024	100	≈59.7	1.676177
768x768	100	≈68.7	1.455097
512x512	100	≈69.4	1.440410
256x256	100	≈69.7	1.434943
128x128	100	≈71.9	1.389953

1.2 Speedup Graph

The graph below visually represents the speedup achieved by the optimization strategies for different image sizes.



1.3 Summary of Results

1024x1024 Image: Achieved the highest performance improvement with a speedup of 1.676177.

768x768 Image: Significant performance improvement with a speedup of 1.455097.

512x512 Image: Noticeable performance improvement with a speedup of 1.440410.

256x256 Image: Performance improvement with a speedup of 1.434943.

128x128 Image: The lowest performance improvement with a speedup of 1.389953.

1.4 Analysis

-Image Size and Performance Improvement of Optimization Strategies

The performance improvement of optimization strategies increases with image size. This can be a

tributed to the fact that larger images benefit more from cache efficiency. Larger images involve more frequent data accesses and exhibit more sequential memory access patterns. Optimization strategies enhance these patterns, improving data locality, which becomes more significant with larger images.

-Data Locality and Cache Misses

Optimization strategies that divide data into smaller blocks make more efficient use of cache space, reducing cache misses and memory access times. A reduction in cache misses leads to shorter CPU memory fetch times, directly improving overall processing time. This effect is particularly pronounced with larger images.

-Specific Advantages of Optimization Strategies

Through optimization, data blocks are loaded into cache more effectively, increasing cache hits during processing. Continuous memory access is facilitated, reducing cache miss penalties. These optimizations reduce CPU idle time and enhance overall processing speed.

The results demonstrate that optimization strategies provide greater performance improvements with larger images. These strategies improve data locality, reduce cache misses, and shorten memory access times, especially for larger images, as evident from the experimental results.

2. Optimization Approach

This report discusses various strategies applied to optimize the baseline implementation of a 3x3 convolution filter in image processing. The goal was to reduce filtering time and enhance performance. Strategies such as blocking, loop unrolling, and memory prefetching were implemented and evaluated for their effectiveness. The baseline and optimized implementations were compared in terms of speedup.

2.1 Optimization Strategies

2.1.1 Blocking

-Strategy: The blocking strategy divides the image into smaller blocks and processes each block individually. This approach enhances data locality, reduces cache misses, and decreases memory access times.

-Implementation: The image is divided into small square blocks, and convolution operations are performed within each block. This allows each block to be effectively loaded into the cache, increasing cache hit rates during processing. For example, the image was divided into 32x32 blocks.

-Evaluation: This strategy achieved a speedup of 1.703361. Blocking effectively improved cache performance and reduced memory access times.

2.1.2 Loop Unrolling

-Strategy: Loop unrolling reduces loop control overhead by increasing the number of operations performed in each iteration. This improves CPU pipeline efficiency and reduces the number of instructions executed.

-Implementation: The inner loop of the convolution operation was unrolled to process multiple pixels in each iteration. This reduces loop control overhead and improves instruction pipeline efficiency. For example, four pixels were processed in each iteration.

-Evaluation: This strategy achieved a speedup of 1.646413. Although it improved performance, it was less effective than blocking due to register pressure and reduced instruction-level parallelism.

2.1.3 Memory Prefetching


-Strategy: Memory prefetching loads data into the cache before it is needed, reducing memory wait times. This technique hides memory access latency and improves performance.

- Implementation:** Functions were used to prefetch data for the next iteration into the cache. This ensures that the data is already in the cache when needed, reducing memory access latency.
- Evaluation:** This strategy achieved a speedup of 1.679224. While data prefetching helped reduce memory access latency, it was not as effective as blocking.

2.1.4 Combining All Optimization Methods

- Strategy:** To maximize performance improvement, all optimization strategies were combined.
- Implementation:** The image was divided into blocks, with loop unrolling and memory prefetching applied within each block. This aimed to leverage the advantages of each optimization strategy.
- Evaluation:** Surprisingly, this combined approach achieved a speedup of 1.590167, the lowest among all strategies. This may be due to several factors:
 - Cache Conflicts: The combined strategies may have increased cache conflicts, reducing performance gains.
 - Instruction Overhead: Managing multiple optimization techniques may have introduced additional instruction overhead.
 - Resource Contention: Combining strategies may have led to contention for CPU resources (e.g., registers, cache lines), reducing efficiency.

2.2 result



Optimization Strategy	Speedup
Blocking	1.703361
Loop Unrolling	1.646413
Memory Prefetching	1.679224
Combined	1.590167

The table and graph clearly demonstrate that the blocking strategy achieved the highest speedup, while the combined approach resulted in the lowest speedup. This indicates that combining multiple optimization strategies requires careful consideration of potential interactions and resource conflicts.

2.3 Analysis

In conclusion, the optimization strategies applied to the baseline implementation showed varying degrees of effectiveness. Blocking was the most effective single strategy, achieving the highest speedup. Loop unrolling and memory prefetching also contributed to performance improvements but were less effective when applied alone. Combining all optimization methods resulted in the lowest speedup, suggesting the need for careful analysis of interactions between optimization strategies. This underscores the importance of evaluating each strategy individually and considering interactions when combining them.